

The Horizon

User Guide — AI Threat Intelligence Dashboard

GETTING STARTED

Logging In


You will receive an email with your login credentials — your registered email address and a temporary password.

1 Open the dashboard link

Go to the URL from the credentials email. You will land on the login page.

2 Sign in with your temporary password

Enter your email address and temporary password, then click **Sign In**.



3 Set your permanent password

On first login you are taken directly to a password setup screen. Create a password meeting all of the following requirements, enter it twice, then click **Set password**. You will be taken to the dashboard.

- At least 8 characters
- At least one uppercase letter
- At least one number
- At least one symbol

If you forget your password, click **Forgot password?** on the login page and a reset link will be sent to your email. To report login issues, contact the CSEC AI Team at [email].

Navigation

After logging in, the dashboard opens to the **Overview** page. The navigation bar at the top provides access to four pages.


PAGE	WHAT IT'S FOR
Overview	AI threat landscape — per-category insights, evidence maturity, trend chart, top threats, and taxonomy coverage
Ask Agent	Natural-language Q&A against the ingested corpus with streamed, cited answers
Sources	Full source table — browse, filter by category and technique, and drill into individual articles
Generate	Download pre-generated briefing slide decks (PPTX) and copy the latest newsletter for email

The 4 Threat Categories

- **Traditional AI Threats** — attacks on ML models, e.g. adversarial examples, data poisoning, model extraction, backdoors
- **LLM Threats** — e.g. prompt injection, jailbreaks, RAG poisoning, guardrail bypass, data leakage
- **Agentic AI Threats** — autonomous agent abuse, e.g. MCP vulnerabilities, tool misuse, coding agent risks
- **AI-Enabled Attacks** — AI as an attack tool, e.g. deepfakes, AI-generated phishing, voice cloning


Overview

The main intelligence view. Shows the AI threat landscape for the selected time window — **Last Week**, **Last Month**, or **Last Quarter**. Use the segmented control at the top right to switch. The ↻ button reloads data; the page also auto-refreshes every 5 minutes.



Category cards

One card per threat category. Click any insight to expand it for the full body, bullet points, and source citations.



Evidence maturity bar

A coloured bar on each card breaks down sources by how well-evidenced they are — from lab research through to threats confirmed in the wild. Click any segment to see the sources behind it.


LEVEL	WHAT IT MEANS
● Research	Demonstrated in papers or controlled lab environments only — no adversary use
● Demonstrated	Working exploit or PoC exists and is reproducible against a real product
● Disclosed	Confirmed vulnerability in a specific product — no working public exploit yet
● Observed	Confirmed in real-world use against real victims — at least one documented incident
● Operational	In sustained or scaled use by one or more threat actors across multiple incidents

Top threats

A ranked list of the highest-importance sources in the period, with maturity badge, publisher, date, and category. Click any title to open the original article.


Taxonomy coverage

A tag matrix showing how many sources are tagged with each specific technique, organised by category. Click any tag to open a panel listing sources for that technique. For a full explanation of all taxonomy codes and technique definitions, go to the Sources page and click the **Taxonomy** button in the top-right.



Ask Agent

A chat interface for querying the ingested corpus in plain language. The agent searches the corpus, retrieves relevant evidence, and streams a structured answer with inline source citations. It searches the last 90 days by default — specify a timeframe in your question for broader or narrower results.



Getting started

On first load, six suggested prompts appear as clickable cards:

- Most important finding
- LLM jailbreak trends
- Agentic AI risks
- MCP vulnerabilities
- AI as an attack tool
- Defender watch list

Click one to send it immediately, or type your own question and press **Enter** or **↑**.

Writing effective questions

Narrow your question by technique, category, actor, or timeframe. Broad questions produce generic answers.

"What prompt injection techniques have been demonstrated against production LLMs in the last three months?"

"Are there documented cases of autonomous AI agents being exploited in the wild?"

"What should defenders watch in the next 90 days?"

Reading the response

Responses are structured with numbered findings, bold key terms, and inline citation numbers. Below each answer, source buttons link to the original articles. A confidence level and temporal scope are shown below the response.

The screenshot shows the 'The Horizon AI THREAT INTELLIGENCE' interface. The top navigation bar includes 'Overview', 'Ask Agent' (active), 'Sources', and 'Generate'. A user profile 'Admin' is in the top right. A blue button at the top right contains the question: 'How is AI being used as an attack tool?'. The response area contains an 'Assessment' paragraph, followed by a numbered finding: '1. Autonomous AI agents are executing complete attack chains end-to-end, without a human operator.' Below this is a summary sentence, followed by two paragraphs of detailed findings with inline citations [3] and [20]. At the bottom, a text input field contains 'Ask about findings, techniques, incidents, trends...' with an upward arrow button. The footer indicates '90-day default · specify timeframe for broader or narrower results' and a 'Clear history' link.

If no relevant sources were found, a **"General knowledge — not grounded in the corpus"** badge appears above the response. Treat these answers with additional caution.

Note


The agent does not search the open web — all answers are drawn from ingested and processed sources.

Very recent articles (last 24–48 hours) may not yet be in the corpus.

Dates reflect when a source was published, not necessarily when the underlying incident occurred — a report published this week may describe an incident that occurred earlier.

Sources

A full table of every ingested and validated source. Use this page to browse the corpus, verify specific articles, or investigate a particular threat area in depth. The **Legend** and **Taxonomy** buttons in the top-right open reference panels explaining all labels, maturity levels, and taxonomy tag codes.



CONTROL

WHAT IT DOES

Time period	Pill switcher (top-right) — filter by 7 days , 30 days , 90 days , or All time
Category tabs	Filter by threat category. Selecting a tab reveals sub-tag pills — click one or more to narrow by specific technique
Label filter	Filter by reading value: Essential , Recommended , Analyst , Background
Importance filter	Filter by evidence maturity: Operational , Observed , Disclosed , Demonstrated , Research
★ Starred	Show only sources you have starred
Reports	Show longer reports and their extracted sub-findings
Sort	Importance (evidence maturity, highest first) or Newest (publish date)

Detail panel

Click any row to expand its detail panel — showing the full summary, evidence maturity and reading value, source type, trust tier, taxonomy tags, and classification rationale. For longer reports it may also include extracted attack walkthroughs, critical insights, and observed trends.

This feature is still being developed and content may vary by source.


Generate

Access pre-generated reports. Two tabs: **Slides** and **Newsletter**.

Slides


Select a reporting period — **1 Month**, **1 Quarter**, or **1 Year** — then click ↓ **Download PPTX**. Each deck includes an executive summary, per-category threat developments, and key insights with source citations.

Slides are AI-generated and must be reviewed and edited by a human before use. Verify claims, check citations, and adjust framing as appropriate before presenting or sharing externally.



Newsletter

Select **Past week** or **Past month** to load the latest newsletter. A live preview renders in the panel. Click **Copy for email** to copy the newsletter to your clipboard, then paste directly into Gmail or Outlook.



If a deck or newsletter shows "Not available yet", it has not been generated for that period. Check back after the next scheduled pipeline run.

Feedback & Issues

This is an early prototype and all feedback is welcome. Contact the CSEC AI Team at **[email]** to report issues or share observations.